

Trees

Dov Kruger

Department of Electrical and Computer Engineering
Rutgers University

January 23, 2024



4.1 Definition of Trees


A tree is an abstract data type representing a hierarchical tree

Consists of nodes, e having 0 or more children

A special node called root is the top of the tree


- Computer science trees grow upside-down!

Nodes with 0 children are called leaf nodes



Binary Trees

In a binary tree all nodes have no more than 2 children




Ordered Binary Trees/Binary Search Tree

An Ordered Binary Tree is a binary tree where


- Each node contains a key
- For each node
 - $\text{key}_{\text{left child}} < \text{key}_{\text{this node}}$
 - $\text{key}_{\text{right child}} \geq \text{key}_{\text{this node}}$or
 - $\text{key}_{\text{left child}} \leq \text{key}_{\text{this node}}$
 - $\text{key}_{\text{right child}} > \text{key}_{\text{this node}}$
- Furthermore this rule applies to the whole subtree

- Example: anything to the right of the root ≥ 6



Inserting into OrderedBinaryTree

```
class Node {  
    public int val;  
    public Node left, right;  
}
```




Tree must define a Node

```
OrderedBinaryTree.insert(v)  
    if root == null  
        root = new Node(v)  
        return // special case for empty tree  
    end
```

R

Insert into OrderedBinaryTree, contd

```
p = root
while p != null
  if p.val < v // can be <= also
    if p.right == null
      p.right <- new Node(v)
      return
    end
  else
    if p.left == null
      p.left = new Node(v)
      return
    end
  end
  p = p.next
```



Recursive Traversal Rules

There are three classic traversal algorithms for binary trees

- Inorder
- Preorder
- Postorder



Inorder

```
inorder(node)
  if node == null
    return // termination condition for recursion!
  end
  inorder(node.left)
  do whatever you want to do to this node (like print it)
  inorder(node.right)
end
```



Preorder

```
preorder(node)
  if node == null
    return // termination condition for recursion!
  end
  do whatever you want to do to this node (like print it)
  preorder(node.left)
  preorder(node.right)
end
```



Postorder

```
postorder(node)
  if node == null
    return // termination condition for recursion!
  end
  postorder(node.left)
  postorder(node.right)
  do whatever you want to do to this node (like print it)
end
```



Another way of Terminating Recursive Traversals

The following code requires testing first that root is not null, but it is more efficient

Complexity is the same

```
postorder(node)
  if node.left != null
    postorder(node.left)
  end
  if node.right != null
    postorder(node.right)
  end
  do whatever you want to do to this node (like print it)
end
```



4.2 Expression Trees


In mathematics we are taught to write expressions inline

$$(x + 2)(x - 5) + 3$$

The operations the above expression represents have an order

Conventions are defined mathematically

The end result is a tree of operations like this:



Three Ways of Writing Expressions

Inorder: $(x + 2)(x - 5)$

Preorder: $* + x 2 - x 5$

Postorder: $x2 + x5 - *$

All 3 reflect the same tree of operations, traversed using inorder, preorder, postorder



Homework: Optimizing an Expression Tree





4.3 Ordered Binary Trees

An ordered binary tree is intrinsically sorted


Just add elements to the tree

Print out the elements using inorder traversal




Deletion from an Ordered Binary Tree

To delete, remove a node, then promote one of the children in its place



Implementation of an Ordered Binary Tree


```
class OrderedBinaryTree {  
private:  
    class Node {  
    public:  
        Node* left;  
        Node* right;  
        int val;  
    };  
    Node* root; // an empty tree has root=null  
    ...  
}
```



Adding a Parent Pointer

For some applications, if the algorithm is looking at a node it is useful to also be able to find the parent

```
class Node {  
public:  
    Node* left;  
    Node* right;  
    Node* parent;  
    int val;  
};
```



Problem with Ordered Binary Trees/Binary Search Trees


Ordered Binary Trees/Binary Search Trees can get lopsided

$O(?)$

- Performance is as bad as lists
- Added overhead of the unused left pointers

Example: Add 1,2,3,4,5 to ordered binary tree


What is the complexity of insertion into a BST?



R

Balanced Case

When a Tree is balanced, n elements fit into $\log_2 n$ levels




$$n = 15 \text{ elements, levels} = \log_2(n + 1) = 4$$




Trees Cannot be Perfectly Balanced

A Tree with one element on a new level is by definition, unbalanced


ok



not ok



resolved



Balanced Binary Trees

Ensure trees remain balanced under all possible operation

Insertion and deletion will have to trigger "rebalancing" while not slowing down



Balanced Binary Trees

There are a number of implementations

- Red-Black Trees (RB-Trees)
- AVL Trees
- Fibonacci Trees

We will just do one: RB-Trees



RB Trees

Objective: Maintain a balanced tree no matter the order of insertion


No worse than $2 \log_2 n$

Rules

- Each node is assigned a color (red or black, really just binary)
- Red nodes cannot have red children

Colors alternate

- Optional: root is black



Definitions for RB-Tree


n_1 .parent is node that has n as a child. This is null for root

n .grandparent() is n 's parent's parent, null if it does not exist

n .uncle is a node's grandparent's other child (not parent)

- If a node has no grandparent, it has no uncle
- n_2 has no uncle

- n_3 uncle is 9



Insertion into RB Tree

Start with an empty tree

```
RBTree.insert(3)
```

```
  calls BST.insert(3)
```

```
  then RBTree.correct(n)
```




$n = \text{root}$

case 1:

```
if n.parent == null
```


```
  return // special case, root
```



Insertion into RB Tree, part 2


```
RBTree.insert(1)  
  calls BST.insert(1)  
  then RBTree.correct(n)
```

```
case 2:  
if parent.color == BLACK  
  return  
end
```




Insertion into RB Tree, part 3

```
RBTree.insert(1)
  calls BST.insert(1)
  then RBTree.correct(n) parent != null
  parent is RED
  case 3:
    if n.parent.isred() and n.uncle().isred()
      flip color of parents
      flip color of grandparent
      RBTree.correct(grandparent)
  recursive
  end
```



Insertion into RB Tree, part 3

```
RBTree.insert(0)  
  calls BST.insert(0)  
  then RBTree.correct(n)  
parent is RED  
[TBD]
```




4.4 BTrees

BTrees are a more complicated generalization of RBTree with a higher degree than 2

Used in database systems to try to keep the depth of the tree to a minimum

Shown: BTree of degree 4



Properties of Disk Access

Hard Disk Drives, and the more modern Solid State Drives (SSD)

- Are Block-oriented devices
- Have high overhead to access a block (latency, and OS call)
- Therefore, try to keep reads to a minimum
- Each block is large, so degree of the BTree can be kept very high (64)

With $n = 64$ (2^6) we can store

Levels	Power of 2	Elements stored
1	2^6	64
2	2^{12}	4096
3	2^{18}	256k
4	2^{24}	16M
5	2^{30}	1 billion



How BTrees are used in Databases

SQL was developed in the 1970s

Creates tables of data, typically with lookup via BTree

The following table defines a primary key

- The records are stored in a BTree using the key
- The records therefore are in sorted order
- Searching for a specific city by name is $O(\log n)$

```
CREATE TABLE Cities (  
  name varchar(10) primary key,  
  population int  
)
```



BTrees in SQL

`SELECT * FROM cities` results sorted by name

`SELECT * FROM cities ORDER BY population` requires sorting

`SELECT * FROM cities WHERE name='New_York'` $O(\log n)$



Adding an Index

`CREATE INDEX cities_bypopulation` Second sorting order
`ON cities (population);`

`SELECT * FROM cities ORDER BY population` no sorting

`SELECT * FROM cities WHERE population=100` $O(\log n)$

`INSERT INTO cities values ('Tokyo', 15000000)` 2x slower



Test Yourself



4.5 Tries

A Trie is a complete tree with a high degree

- 26 for lowercase English letters
- Has not been used much because storage requirements are high
- Hash maps have dominated tries for most applications

Today, there are applications

- Autocompletion of words
- DNA searches

<https://brilliant.org/wiki/tries/>



Code Structure of a Trie

A

Trie node has

- boolean stating whether the path is a word or not
- Pointers to the next node


A trie dictionary can contain the root node

- Does not have to be a pointer
- Can never be null

```
class Trie {  
    struct Node {  
        bool isWord;  
        Node* next[26];  
    }  
    Node root;  
public:  
    void add(const string& w)  
    void remove(const string& w)  
    bool contains(const string& w)  
    bool startsWith(const string& w)  
};
```




Inserting into a Trie




R

Searching for a Word in a Trie



R

Searching for a Prefix in a Trie



R

The game of Boggle Excellent application using tries.
We regularly assign this for homework. Write a program that searches the board

